

# Homework 3: Multi-Agent Search

Please keep the title of each section and delete examples.

## Part I. Implementation (5%):

### Part1

```
136 # Begin your code (Part 1)
137 """
138 Using recursion to compute the minimax values of each state. The minimax function takes state, depth, agent index as arguments,
139 and returns minimax value of the state. If the state is a terminal state or the maximum depth has been reached, the function
140 returns the evaluation of the state. Otherwise, the function computes the minimax value for the current agent by either
141 maximizing or minimizing the values of the next agent's actions, depending on the current agent is Pacman or a ghost.
142 Finally, the function returns the computed minimax value.
143 """
144 def minimax(state, depth, agentIndex):
145     if state.isWin() or state.islose() or depth == self.depth:
146         return self.evaluationFunction(state)
147     if agentIndex == 0: # maximize for pacman
148         v = float('-inf')
149         for action in state.getLegalActions(agentIndex):
150             v = max(v, minimax(state.getNextState(agentIndex, action), depth, agentIndex + 1))
151         return v
152     else: # minimize for ghost
153         v = float('inf')
154         for action in state.getLegalActions(agentIndex):
155             if agentIndex == state.getNumAgents() - 1: # next agent is pacman, depth + 1
156                 v = min(v, minimax(state.getNextState(agentIndex, action), depth + 1, 0))
157             else: # move to next ghost
158                 v = min(v, minimax(state.getNextState(agentIndex, action), depth, agentIndex + 1))
159         return v
160 legalMoves = gameState.getLegalActions() # get all legal actions
161 # get the score for each legal action by running minimax on the resulting state
162 scores = [minimax(gameState.getNextState(0, action), 0, 1) for action in legalMoves]
163 # find the best score and choose a random action with the best score
164 bestScore = max(scores)
165 bestIndices = [index for index in range(len(scores)) if scores[index] == bestScore]
166 chosenIndex = random.choice(bestIndices)
167 return legalMoves[chosenIndex]
168 # End your code (Part 1)
```

## Part2:

```
180 # Begin your code (Part 2)
181 """
182 maxValue function returns the maximum value achievable by the current agent in the current state,
183 while minValue function returns the minimum value achievable by the next agent.
184 The main loop iterates through all legal actions of the current agent and calculates the value
185 of each action using the minValue function, and updates the best action accordingly. It uses
186 alpha-beta pruning to avoid exploring paths that will not lead to a better outcome.
187 Finally, return the best action.
188 """
189 def maxValue(state, agentIndex, depth, alpha, beta):
190     v = float("-inf")
191     legalActions = state.getLegalActions(agentIndex)
192     if not legalActions or depth == self.depth:
193         return self.evaluationFunction(state)
194     for action in legalActions:
195         nextState = state.getNextState(agentIndex, action)
196         v = max(v, minValue(nextState, agentIndex + 1, depth, alpha, beta))
197         if v > beta:
198             return v
199         alpha = max(alpha, v)
200     return v
201
202 def minValue(state, agentIndex, depth, alpha, beta):
203     v = float("inf")
204     legalActions = state.getLegalActions(agentIndex)
205     if not legalActions or depth == self.depth:
206         return self.evaluationFunction(state)
207     for action in legalActions:
208         nextState = state.getNextState(agentIndex, action)
209         if agentIndex == state.getNumAgents() - 1:
210             v = min(v, maxValue(nextState, 0, depth + 1, alpha, beta))
211         else:
212             v = min(v, minValue(nextState, agentIndex + 1, depth, alpha, beta))
213         if v < alpha:
214             return v
215         beta = min(beta, v)
216     return v
217
218 legalActions = gameState.getLegalActions()
219 bestAction = None
220 v = float("-inf")
221 alpha = float("-inf")
222 beta = float("inf")
223 for action in legalActions:
224     nextState = gameState.getNextState(0, action)
225     nextValue = minValue(nextState, 1, 0, alpha, beta)
226     if nextValue > v:
227         v = nextValue
228         bestAction = action
229     if v > beta:
230         return bestAction
231     alpha = max(alpha, v)
232 return bestAction
233 # End your code (Part 2)
```

### Part3:

```
248 # Begin your code (Part 3)
249 """
250 maxValue function recursively calculates the maximum value of a given state by iterating over
251 all legal actions of the pacman player and calling the expectValue function on the resulting next state.
252 The expectValue function calculates the expected value of a given state by iterating over all legal
253 actions of the current agent and recursively calling either maxValue or expectValue on the resulting next state.
254 And returns the average of the resulting values, weighted by probability of each action.
255 Finally, the agent selects the action with the highest expected score and returns.
256 """
257
258 def maxValue(state, depth):
259     if state.isWin() or state.isLose() or depth == self.depth:
260         return self.evaluationFunction(state)
261     v = float('-inf')
262     for action in state.getLegalActions():
263         v = max(v, expectValue(state.getNextState(0, action), depth, 1))
264     return v
265
266 def expectValue(state, depth, agentIndex):
267     if state.isWin() or state.isLose() or depth == self.depth:
268         return self.evaluationFunction(state)
269     v = 0
270     legalActions = state.getLegalActions(agentIndex)
271     p = 1.0 / len(legalActions)
272     for action in legalActions:
273         if agentIndex == state.getNumAgents() - 1:
274             v += p * maxValue(state.getNextState(agentIndex, action), depth + 1)
275         else:
276             v += p * expectValue(state.getNextState(agentIndex, action), depth, agentIndex + 1)
277     return v
278
279 legalMoves = gameState.getLegalActions()
280 scores = []
281 for action in legalMoves:
282     score = expectValue(gameState.getNextState(0, action), 0, 1)
283     scores.append(score)
284
285 bestScore = max(scores)
286 bestIndices = [index for index in range(len(scores)) if scores[index] == bestScore]
287 chosenIndex = random.choice(bestIndices)
288 return legalMoves[chosenIndex]
289 # End your code (Part 3)
```

### Part4:

```
295 # Begin your code (Part 4)
296 """
297 The function calculates a score for a pacman game state based on the distance and proximity of ghosts to pacman,
298 the distance to the closest remaining food, and the number of remaining capsules.
299 The function prioritizes getting closer to food while avoiding ghosts, while also
300 taking into account the remaining capsules on the board.
301 """
302
303 pPos = currentGameState.getPacmanPosition()
304 gDist = 0
305 gProximity = 0
306 for gState in currentGameState.getGhostPositions():
307     dist = util.manhattanDistance(pPos, gState)
308     gDist += dist
309     if dist <= 1:
310         gProximity += 1
311
312 foodList = currentGameState.getFood().asList()
313 minfoodList = min([util.manhattanDistance(pPos, foodPos) for foodPos in foodList], default=0)
314
315 capsuleNum = len(currentGameState.getCapsules())
316 return currentGameState.getScore() - (1 / (gDist+1)) - gProximity + (1 / (minfoodList+1)) - capsuleNum
317 # End your code (Part 4)
```

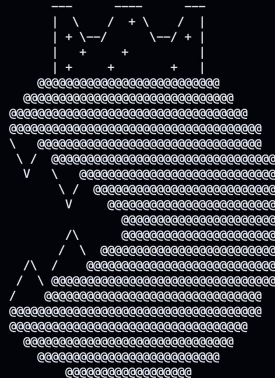
## Part II. Results & Analysis (5%):

### Provisional grades

Question part1: 20/20  
Question part2: 25/25  
Question part3: 25/25  
Question part4: 10/10

Total: 80/80

ALL HAIL GRANDPAC.  
LONG LIVE THE GHOSTBUSTING KING.



### Question part4

```
=====  
Pacman emerges victorious! Score: 1373  
Pacman emerges victorious! Score: 1159  
Pacman emerges victorious! Score: 1158  
Pacman emerges victorious! Score: 1174  
Pacman emerges victorious! Score: 1371  
Pacman emerges victorious! Score: 849  
Pacman emerges victorious! Score: 1170  
Pacman emerges victorious! Score: 926  
Pacman emerges victorious! Score: 1155  
Pacman emerges victorious! Score: 1160  
Average Score: 1149.5  
Scores: 1373.0, 1159.0, 1158.0, 1174.0, 1371.0, 849.0, 1170.0, 926.0, 1155.0, 1160.0  
Win Rate: 10/10 (1.00)  
Record: Win, Win, Win, Win, Win, Win, Win, Win, Win, Win  
*** PASS: test_cases/part4/grade-agent.test (8 of 8 points)  
*** EXTRA CREDIT: 2 points  
*** 1149.5 average score (4 of 4 points)  
*** Grading scheme:  
*** < 500: 0 points  
*** >= 500: 2 points  
*** >= 1000: 4 points  
*** 10 games not timed out (2 of 2 points)  
*** Grading scheme:  
*** < 0: fail  
*** >= 0: 0 points  
*** >= 5: 1 points  
*** >= 10: 2 points  
*** 10 wins (4 of 4 points)  
*** Grading scheme:  
*** < 1: fail  
*** >= 1: 1 points  
*** >= 4: 2 points  
*** >= 7: 3 points  
*** >= 10: 4 points  
### Question part4: 10/10 ###
```

During the evaluation of the pacman game state, the distance between pacman and ghosts, food, the proximity of ghosts around pacman, and the number of capsules are taken into consideration. Among these factors, the distance between pacman and food has the most significant impact on the score. This makes sense since the ultimate objective of the game is to gain more points, and getting closer to food is the key to achieving this objective. However, the other factors such as the proximity of ghosts and the number of capsules also influence the outcome since they are essential for survival and successful completion of the game.